

softmax

September 29, 2021

1 Softmax exercise

Complete and hand in this completed worksheet (including its outputs and any supporting code outside of the worksheet) with your assignment submission. For more details see the [assignments page](#) on the course website.

This exercise is analogous to the SVM exercise. You will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```
[1]: # Setup Colab connection with Drive
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)

# Define working directory
FOLDERNAME = 'Coursework/Fall 2021/682/Assignments/assignment1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

import sys
sys.path.append(f'/content/gdrive/My Drive/{FOLDERNAME}')
```

Mounted at /content/gdrive

```
[2]: from __future__ import print_function
import random
import numpy as np
from cs682.data_utils import load_CIFAR10
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
```

```
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading extenrnal modules
# see http://stackoverflow.com/questions/1907993/
→ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

```
[3]: #Switch to working directory
%cd gdrive/My\ Drive/$FOLDERNAME
```

/content/gdrive/My Drive/Coursework/Fall 2021/682/Assignments/assignment1

```
[4]: def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000,
→ num_dev=500):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the linear classifier. These are the same steps as we used for the
    SVM, but condensed to a single function.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs682/datasets/cifar-10-batches-py'

    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # subsample the data
    mask = list(range(num_training, num_training + num_validation))
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = list(range(num_training))
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = list(range(num_test))
    X_test = X_test[mask]
    y_test = y_test[mask]
    mask = np.random.choice(num_training, num_dev, replace=False)
    X_dev = X_train[mask]
    y_dev = y_train[mask]

    # Preprocessing: reshape the image data into rows
    X_train = np.reshape(X_train, (X_train.shape[0], -1))
    X_val = np.reshape(X_val, (X_val.shape[0], -1))
    X_test = np.reshape(X_test, (X_test.shape[0], -1))
    X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis = 0)
    X_train -= mean_image
```

```

X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image

# add bias dimension and transform into columns
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev

# Cleaning up variables to prevent loading data multiple times (which may cause
→memory issue)
try:
    del X_train, y_train
    del X_test, y_test
    print('Clear previously loaded data.')
except:
    pass

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev =
    →get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
print('dev data shape: ', X_dev.shape)
print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

1.1 Softmax Classifier

Your code for this section will all be written inside `cs682/classifiers/softmax.py`.

```
[5]: # First implement the naive softmax loss function with nested loops.
# Open the file cs682/classifiers/softmax.py and implement the
# softmax_loss_naive function.

from cs682.classifiers.softmax import softmax_loss_naive
import time

# Generate a random softmax weight matrix and use it to compute the loss.
W = np.random.randn(3073, 10) * 0.0001
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As a rough sanity check, our loss should be something close to -log(0.1).
print('loss: %f' % loss)
print('sanity check: %f' % (-np.log(0.1)))
```

loss: 2.323482

sanity check: 2.302585

1.2 Inline Question 1:

Why do we expect our loss to be close to $-\log(0.1)$? Explain briefly.**

Your answer: We expect loss to be close to $-\log(0.1)$ as we are not learning any parameters at this point and we expect each class (out of 10) to be equally likely when using a randomly initialized initial W.

```
[6]: # Complete the implementation of softmax_loss_naive and implement a (naive)
# version of the gradient that uses nested loops.
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 0.0)

# As we did for the SVM, use numeric gradient checking as a debugging tool.
# The numeric gradient should be close to the analytic gradient.
from cs682.gradient_check import grad_check_sparse
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)

# similar to SVM case, do another gradient check with regularization
loss, grad = softmax_loss_naive(W, X_dev, y_dev, 5e1)
f = lambda w: softmax_loss_naive(w, X_dev, y_dev, 5e1)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

numerical: -0.284203 analytic: -0.284203, relative error: 8.187001e-08

numerical: 0.725570 analytic: 0.725570, relative error: 1.115359e-07

numerical: 0.915656 analytic: 0.915656, relative error: 4.604942e-08

numerical: 3.544753 analytic: 3.544753, relative error: 1.267183e-08

numerical: 0.609784 analytic: 0.609784, relative error: 1.791180e-07

numerical: 0.472553 analytic: 0.472553, relative error: 1.343816e-07

numerical: 1.911679 analytic: 1.911679, relative error: 5.556583e-08

numerical: -0.878860 analytic: -0.878860, relative error: 3.859701e-08

```

numerical: 1.211142 analytic: 1.211142, relative error: 5.585657e-08
numerical: 1.645632 analytic: 1.645632, relative error: 2.717919e-08
numerical: 0.447980 analytic: 0.447980, relative error: 7.554675e-08
numerical: -1.055898 analytic: -1.055898, relative error: 3.890911e-10
numerical: -0.960308 analytic: -0.960308, relative error: 1.570917e-08
numerical: -2.028626 analytic: -2.028626, relative error: 9.224981e-09
numerical: 2.628991 analytic: 2.628991, relative error: 2.425265e-08
numerical: 2.341801 analytic: 2.341801, relative error: 1.486300e-08
numerical: 0.507949 analytic: 0.507949, relative error: 1.452742e-07
numerical: 2.391970 analytic: 2.391970, relative error: 1.223744e-08
numerical: -2.085640 analytic: -2.085640, relative error: 2.415905e-08
numerical: -1.401349 analytic: -1.401349, relative error: 3.581999e-08

```

```

[7]: # Now that we have a naive implementation of the softmax loss function and its
      ↪gradient,
      # implement a vectorized version in softmax_loss_vectorized.
      # The two versions should compute the same results, but the vectorized version
      ↪should be
      # much faster.
      tic = time.time()
      loss_naive, grad_naive = softmax_loss_naive(W, X_dev, y_dev, 0.000005)
      toc = time.time()
      print('naive loss: %e computed in %fs' % (loss_naive, toc - tic))

      from cs682.classifiers.softmax import softmax_loss_vectorized
      tic = time.time()
      loss_vectorized, grad_vectorized = softmax_loss_vectorized(W, X_dev, y_dev, 0.
      ↪000005)
      toc = time.time()
      print('vectorized loss: %e computed in %fs' % (loss_vectorized, toc - tic))

      # As we did for the SVM, we use the Frobenius norm to compare the two versions
      # of the gradient.
      grad_difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
      print('Loss difference: %f' % np.abs(loss_naive - loss_vectorized))
      print('Gradient difference: %f' % grad_difference)

```

```

naive loss: 2.323482e+00 computed in 0.113843s
vectorized loss: 2.323482e+00 computed in 0.017536s
Loss difference: 0.000000
Gradient difference: 0.000000

```

```

[8]: # Use the validation set to tune hyperparameters (regularization strength and
      # learning rate). You should experiment with different ranges for the learning
      # rates and regularization strengths; if you are careful you should be able to
      # get a classification accuracy of over 0.35 on the validation set.
      from cs682.classifiers import Softmax

```

```

results = {}
best_val = -1
best_softmax = None
learning_rates = [1e-7, 2e-7, 3e-7, 4e-7, 5e-7]
regularization_strengths = [2.5e4, 3e4, 3.5e4, 4e4, 4.5e4, 5e4]

#####
# TODO:
# Use the validation set to set the learning rate and regularization strength.
# This should be identical to the validation that you did for the SVM; save
# the best trained softmax classifier in best_softmax.
# Your code
with np.errstate(all='ignore'):
    for lr in learning_rates:
        for r in regularization_strengths:
            sm = Softmax()
            loss_hist = sm.train(X_train, y_train, learning_rate=lr, reg=r,
num_iters=1500, verbose=False)
            y_train_pred = sm.predict(X_train)
            train_acc = np.mean(y_train == y_train_pred)
            y_val_pred = sm.predict(X_val)
            val_acc = np.mean(y_val == y_val_pred)
            results[(lr,r)] = (train_acc, val_acc)
            if val_acc > best_val:
                best_val = val_acc
                best_softmax = sm
#####
#
#
#
# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))

print('best validation accuracy achieved during cross-validation: %f' %
best_val)

```

```

lr 1.000000e-07 reg 2.500000e+04 train accuracy: 0.335592 val accuracy: 0.347000
lr 1.000000e-07 reg 3.000000e+04 train accuracy: 0.324735 val accuracy: 0.333000

```

```

lr 1.000000e-07 reg 3.500000e+04 train accuracy: 0.316980 val accuracy: 0.336000
lr 1.000000e-07 reg 4.000000e+04 train accuracy: 0.317735 val accuracy: 0.329000
lr 1.000000e-07 reg 4.500000e+04 train accuracy: 0.310122 val accuracy: 0.326000
lr 1.000000e-07 reg 5.000000e+04 train accuracy: 0.308143 val accuracy: 0.318000
lr 2.000000e-07 reg 2.500000e+04 train accuracy: 0.329653 val accuracy: 0.343000
lr 2.000000e-07 reg 3.000000e+04 train accuracy: 0.318816 val accuracy: 0.339000
lr 2.000000e-07 reg 3.500000e+04 train accuracy: 0.323020 val accuracy: 0.336000
lr 2.000000e-07 reg 4.000000e+04 train accuracy: 0.318163 val accuracy: 0.336000
lr 2.000000e-07 reg 4.500000e+04 train accuracy: 0.304531 val accuracy: 0.320000
lr 2.000000e-07 reg 5.000000e+04 train accuracy: 0.293571 val accuracy: 0.316000
lr 3.000000e-07 reg 2.500000e+04 train accuracy: 0.333571 val accuracy: 0.343000
lr 3.000000e-07 reg 3.000000e+04 train accuracy: 0.316816 val accuracy: 0.333000
lr 3.000000e-07 reg 3.500000e+04 train accuracy: 0.328102 val accuracy: 0.340000
lr 3.000000e-07 reg 4.000000e+04 train accuracy: 0.311367 val accuracy: 0.329000
lr 3.000000e-07 reg 4.500000e+04 train accuracy: 0.317673 val accuracy: 0.323000
lr 3.000000e-07 reg 5.000000e+04 train accuracy: 0.298061 val accuracy: 0.308000
lr 4.000000e-07 reg 2.500000e+04 train accuracy: 0.328469 val accuracy: 0.343000
lr 4.000000e-07 reg 3.000000e+04 train accuracy: 0.321224 val accuracy: 0.338000
lr 4.000000e-07 reg 3.500000e+04 train accuracy: 0.308816 val accuracy: 0.328000
lr 4.000000e-07 reg 4.000000e+04 train accuracy: 0.314102 val accuracy: 0.329000
lr 4.000000e-07 reg 4.500000e+04 train accuracy: 0.310143 val accuracy: 0.324000
lr 4.000000e-07 reg 5.000000e+04 train accuracy: 0.305918 val accuracy: 0.310000
lr 5.000000e-07 reg 2.500000e+04 train accuracy: 0.324776 val accuracy: 0.330000
lr 5.000000e-07 reg 3.000000e+04 train accuracy: 0.322633 val accuracy: 0.330000
lr 5.000000e-07 reg 3.500000e+04 train accuracy: 0.318041 val accuracy: 0.334000
lr 5.000000e-07 reg 4.000000e+04 train accuracy: 0.311837 val accuracy: 0.333000
lr 5.000000e-07 reg 4.500000e+04 train accuracy: 0.312224 val accuracy: 0.328000
lr 5.000000e-07 reg 5.000000e+04 train accuracy: 0.307531 val accuracy: 0.321000
best validation accuracy achieved during cross-validation: 0.347000

```

```

[9]: # evaluate on test set
# Evaluate the best softmax on test set
y_test_pred = best_softmax.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))

```

softmax on raw pixels final test set accuracy: 0.343000

Inline Question - True or False

It's possible to add a new datapoint to a training set that would leave the SVM loss unchanged, but this is not the case with the Softmax classifier loss.

Your answer: True

Your explanation: It is possible to add a data point which is outside the SVM margin, and hence would contribute 0 to the loss. However addition of a data point would change the data probability distribution and hence would affect the softmax loss.

```

[10]: # Visualize the learned weights for each class
w = best_softmax.W[:-1,:] # strip out the bias

```

```

w = w.reshape(32, 32, 3, 10)

w_min, w_max = np.min(w), np.max(w)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype('uint8'))
    plt.axis('off')
    plt.title(classes[i])

```

