

# PyTorch

October 30, 2021

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/gdrive', force_remount=True)

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cs231n/assignments/assignment1/'
FOLDERNAME = 'Coursework/Fall 2021/682/Assignments/assignment2'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/gdrive/My Drive/{}'.format(FOLDERNAME))

#Switch to working directory
%cd /content/gdrive/My\ Drive/$FOLDERNAME
```

```
Mounted at /content/gdrive
/content/gdrive/My Drive/Coursework/Fall 2021/682/Assignments/assignment2
```

## 1 What's this PyTorch business?

You've written a lot of code in this assignment to provide a whole host of neural network functionality. Dropout, Batch Norm, and 2D convolutions are some of the workhorses of deep learning in computer vision. You've also worked hard to make your code efficient and vectorized.

For the last part of this assignment, though, we're going to leave behind your beautiful code-base and instead migrate to one of two popular deep learning frameworks: in this instance, PyTorch (or TensorFlow, if you switch over to that notebook).

### 1.0.1 What is PyTorch?

PyTorch is a system for executing dynamic computational graphs over Tensor objects that behave similarly as numpy ndarray. It comes with a powerful automatic differentiation engine that removes the need for manual back-propagation.

### 1.0.2 Why do we use deep learning frameworks?

- Our code will now run on GPUs! Much faster training. When using a framework like PyTorch or TensorFlow you can harness the power of the GPU for your own custom neural network architectures without having to write CUDA code directly (which is beyond the scope of this class).
- We want you to be ready to use one of these frameworks for your project so you can experiment more efficiently than if you were writing every feature you want to use by hand.
- We want you to stand on the shoulders of giants! TensorFlow and PyTorch are both excellent frameworks that will make your lives a lot easier, and now that you understand their guts, you are free to use them :)
- We want you to be exposed to the sort of deep learning code you might run into in academia or industry.

### 1.1 How will I learn PyTorch?

Justin Johnson has made an excellent [tutorial](#) for PyTorch.

You can also find the detailed [API doc](#) here. If you have other questions that are not addressed by the API docs, the [PyTorch forum](#) is a much better place to ask than StackOverflow.

## 2 Table of Contents

This assignment has 5 parts. You will learn PyTorch on different levels of abstractions, which will help you understand it better and prepare you for the final project.

1. Preparation: we will use CIFAR-10 dataset.
2. Barebones PyTorch: we will work directly with the lowest-level PyTorch Tensors.
3. PyTorch Module API: we will use `nn.Module` to define arbitrary neural network architecture.
4. PyTorch Sequential API: we will use `nn.Sequential` to define a linear feed-forward network very conveniently.
5. CIFAR-10 open-ended challenge: please implement your own network to get as high accuracy as possible on CIFAR-10. You can experiment with any layer, optimizer, hyperparameters or other advanced features.

Here is a table of comparison:

API	Flexibility	Convenience
Barebone	High	Low
<code>nn.Module</code>	High	Medium
<code>nn.Sequential</code>	Low	High

## 3 GPU

You can manually switch to a GPU device on Colab by clicking Runtime -> Change runtime type and selecting GPU under Hardware Accelerator. You should do this before running the following cells to import packages, since the kernel gets restarted upon switching runtimes.

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader
from torch.utils.data import sampler

import torchvision.datasets as dset
import torchvision.transforms as T

import numpy as np
```

```
[2]: USE_GPU = True

dtype = torch.float32 # we will be using float throughout this tutorial

if USE_GPU and torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')

# Constant to control how frequently we print train loss
print_every = 100

print('using device:', device)
```

using device: cuda

## 4 Part I. Preparation

First, we load the CIFAR-10 dataset. This might take a couple minutes the first time you do it, but the files should stay cached after that.

In previous parts of the assignment we had to write our own code to download the CIFAR-10 dataset, preprocess it, and iterate through it in minibatches; PyTorch provides convenient tools to automate this process for us.

```
[43]: NUM_TRAIN = 49000

# The torchvision.transforms package provides tools for preprocessing data
# and for performing data augmentation; here we set up a transform to
# preprocess the data by subtracting the mean RGB value and dividing by the
# standard deviation of each RGB value; we've hardcoded the mean and std.
transform = T.Compose([
    T.ToTensor(),
    T.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010))
])

# We set up a Dataset object for each split (train / val / test); Datasets load
```

```

# training examples one at a time, so we wrap each Dataset in a DataLoader,
→which
# iterates through the Dataset and forms minibatches. We divide the CIFAR-10
# training set into train and val sets by passing a Sampler object to the
# DataLoader telling how it should sample from the underlying Dataset.
cifar10_train = dset.CIFAR10('./cs682/datasets', train=True, download=True,
                             transform=transform)
loader_train = DataLoader(cifar10_train, batch_size=64,
                          sampler=sampler.SubsetRandomSampler(range(NUM_TRAIN)))

cifar10_val = dset.CIFAR10('./cs682/datasets', train=True, download=True,
                           transform=transform)
loader_val = DataLoader(cifar10_val, batch_size=64,
                       sampler=sampler.SubsetRandomSampler(range(NUM_TRAIN,
→50000))))

cifar10_test = dset.CIFAR10('./cs682/datasets', train=False, download=True,
                            transform=transform)
loader_test = DataLoader(cifar10_test, batch_size=64)

```

Files already downloaded and verified

Files already downloaded and verified

Files already downloaded and verified

You have an option to use **GPU** by setting the flag to **True** below. It is not necessary to use GPU for this assignment. Note that if your computer does not have CUDA enabled, `torch.cuda.is_available()` will return `False` and this notebook will fallback to CPU mode.

The global variables `dtype` and `device` will control the data types throughout this assignment.

## 5 Part II. Barebones PyTorch

PyTorch ships with high-level APIs to help us define model architectures conveniently, which we will cover in Part II of this tutorial. In this section, we will start with the barebone PyTorch elements to understand the autograd engine better. After this exercise, you will come to appreciate the high-level model API more.

We will start with a simple fully-connected ReLU network with two hidden layers and no biases for CIFAR classification. This implementation computes the forward pass using operations on PyTorch Tensors, and uses PyTorch autograd to compute gradients. It is important that you understand every line, because you will write a harder version after the example.

When we create a PyTorch Tensor with `requires_grad=True`, then operations involving that Tensor will not just compute values; they will also build up a computational graph in the background, allowing us to easily backpropagate through the graph to compute gradients of some Tensors with respect to a downstream loss. Concretely if `x` is a Tensor with `x.requires_grad == True` then after backpropagation `x.grad` will be another Tensor holding the gradient of `x` with respect to the scalar loss at the end.

### 5.0.1 PyTorch Tensors: Flatten Function

A PyTorch Tensor is conceptionally similar to a numpy array: it is an n-dimensional grid of numbers, and like numpy PyTorch provides many functions to efficiently operate on Tensors. As a simple example, we provide a `flatten` function below which reshapes image data for use in a fully-connected neural network.

Recall that image data is typically stored in a Tensor of shape  $N \times C \times H \times W$ , where:

- $N$  is the number of datapoints
- $C$  is the number of channels
- $H$  is the height of the intermediate feature map in pixels
- $W$  is the width of the intermediate feature map in pixels

This is the right way to represent the data when we are doing something like a 2D convolution, that needs spatial understanding of where the intermediate features are relative to each other. When we use fully connected affine layers to process the image, however, we want each datapoint to be represented by a single vector -- it's no longer useful to segregate the different channels, rows, and columns of the data. So, we use a "flatten" operation to collapse the  $C \times H \times W$  values per representation into a single long vector. The `flatten` function below first reads in the  $N$ ,  $C$ ,  $H$ , and  $W$  values from a given batch of data, and then returns a "view" of that data. "View" is analogous to numpy's "reshape" method: it reshapes  $x$ 's dimensions to be  $N \times ??$ , where  $??$  is allowed to be anything (in this case, it will be  $C \times H \times W$ , but we don't need to specify that explicitly).

```
[44]: def flatten(x):
      N = x.shape[0] # read in N, C, H, W
      return x.view(N, -1) # "flatten" the C * H * W values into a single vector
      ↪per image

def test_flatten():
    x = torch.arange(12).view(2, 1, 3, 2)
    print('Before flattening: ', x)
    print('After flattening: ', flatten(x))
    print(x.shape)
    print(flatten(x).shape)

test_flatten()
```

```
Before flattening: tensor([[[[ 0,  1],
      [ 2,  3],
      [ 4,  5]]],
```

```
      [[[ 6,  7],
      [ 8,  9],
      [10, 11]]]])
```

```
After flattening: tensor([[ 0,  1,  2,  3,  4,  5],
      [ 6,  7,  8,  9, 10, 11]])
```

```
torch.Size([2, 1, 3, 2])
```

```
torch.Size([2, 6])
```

### 5.0.2 Barebones PyTorch: Two-Layer Network

Here we define a function `two_layer_fc` which performs the forward pass of a two-layer fully-connected ReLU network on a batch of image data. After defining the forward pass we check that it doesn't crash and that it produces outputs of the right shape by running zeros through the network.

You don't have to write any code here, but it's important that you read and understand the implementation.

```
[45]: import torch.nn.functional as F # useful stateless functions

def two_layer_fc(x, params):
    """
    A fully-connected neural networks; the architecture is:
    NN is fully connected -> ReLU -> fully connected layer.
    Note that this function only defines the forward pass;
    PyTorch will take care of the backward pass for us.

    The input to the network will be a minibatch of data, of shape
    (N, d1, ..., dM) where  $d1 * \dots * dM = D$ . The hidden layer will have  $H$ 
    units,
    and the output layer will produce scores for  $C$  classes.

    Inputs:
    - x: A PyTorch Tensor of shape (N, d1, ..., dM) giving a minibatch of
      input data.
    - params: A list [w1, w2] of PyTorch Tensors giving weights for the network;
      w1 has shape (D, H) and w2 has shape (H, C).

    Returns:
    - scores: A PyTorch Tensor of shape (N, C) giving classification scores for
      the input data x.
    """
    # first we flatten the image
    x = flatten(x) # shape: [batch_size, C x H x W]

    w1, w2 = params

    # Forward pass: compute predicted y using operations on Tensors. Since w1
    and
    # w2 have requires_grad=True, operations involving these Tensors will cause
    # PyTorch to build a computational graph, allowing automatic computation of
    # gradients. Since we are no longer implementing the backward pass by hand
    we
    # don't need to keep references to intermediate values.
    # you can also use `.clamp(min=0)`, equivalent to F.relu()
    x = F.relu(x.mm(w1))
    x = x.mm(w2)
```

```

    return x

def two_layer_fc_test():
    hidden_layer_size = 42
    x = torch.zeros((64, 50), dtype=dtype) # minibatch size 64, feature
    →dimension 50
    w1 = torch.zeros((50, hidden_layer_size), dtype=dtype)
    w2 = torch.zeros((hidden_layer_size, 10), dtype=dtype)
    scores = two_layer_fc(x, [w1, w2])
    print(scores.size()) # you should see [64, 10]

two_layer_fc_test()

```

torch.Size([64, 10])

### 5.0.3 Barebones PyTorch: Three-Layer ConvNet

Here you will complete the implementation of the function `three_layer_convnet`, which will perform the forward pass of a three-layer convolutional network. Like above, we can immediately test our implementation by passing zeros through the network. The network should have the following architecture:

1. A convolutional layer (with bias) with `channel_1` filters, each with shape `KW1 x KH1`, and zero-padding of two
2. ReLU nonlinearity
3. A convolutional layer (with bias) with `channel_2` filters, each with shape `KW2 x KH2`, and zero-padding of one
4. ReLU nonlinearity
5. Fully-connected layer with bias, producing scores for `C` classes.

**HINT:** For convolutions: <http://pytorch.org/docs/stable/nn.html#torch.nn.functional.conv2d>; pay attention to the shapes of convolutional filters!

```

[47]: def three_layer_convnet(x, params):
    """
    Performs the forward pass of a three-layer convolutional network with the
    architecture defined above.

    Inputs:
    - x: A PyTorch Tensor of shape (N, 3, H, W) giving a minibatch of images
    - params: A list of PyTorch Tensors giving the weights and biases for the
      network; should contain the following:
      - conv_w1: PyTorch Tensor of shape (channel_1, 3, KH1, KW1) giving
    →weights
        for the first convolutional layer
      - conv_b1: PyTorch Tensor of shape (channel_1,) giving biases for the
    →first

```

```

        convolutional layer
        - conv_w2: PyTorch Tensor of shape (channel_2, channel_1, KH2, KW2)
→giving
        weights for the second convolutional layer
        - conv_b2: PyTorch Tensor of shape (channel_2,) giving biases for the
→second
        convolutional layer
        - fc_w: PyTorch Tensor giving weights for the fully-connected layer. Can
→you
        figure out what the shape should be?
        - fc_b: PyTorch Tensor giving biases for the fully-connected layer. Can
→you
        figure out what the shape should be?

Returns:
- scores: PyTorch Tensor of shape (N, C) giving classification scores for x
"""
conv_w1, conv_b1, conv_w2, conv_b2, fc_w, fc_b = params
scores = None

→#####
# TODO: Implement the forward pass for the three-layer ConvNet.
→#

→#####
x = F.conv2d(x, conv_w1, conv_b1, stride=1, padding=2)
x = F.relu(x)
x = F.conv2d(x, conv_w2, conv_b2, stride=1, padding=1)
x = F.relu(x)
x = flatten(x)
x = x.mm(fc_w) + fc_b
scores = x
pass

→#####
#                               END OF YOUR CODE
→#

→#####
return scores

```

After defining the forward pass of the ConvNet above, run the following cell to test your implementation.

When you run this function, scores should have shape (64, 10).

```
[48]: def three_layer_convnet_test():
```



```

x = torch.zeros((64, 3, 32, 32), dtype=dtype) # minibatch size 64, image
→size [3, 32, 32]

conv_w1 = torch.zeros((6, 3, 5, 5), dtype=dtype) # [out_channel,
→in_channel, kernel_H, kernel_W]
conv_b1 = torch.zeros((6,)) # out_channel
conv_w2 = torch.zeros((9, 6, 3, 3), dtype=dtype) # [out_channel,
→in_channel, kernel_H, kernel_W]
conv_b2 = torch.zeros((9,)) # out_channel

# you must calculate the shape of the tensor after two conv layers, before
→the fully-connected layer
fc_w = torch.zeros((9 * 32 * 32, 10))
fc_b = torch.zeros(10)

scores = three_layer_convnet(x, [conv_w1, conv_b1, conv_w2, conv_b2, fc_w,
→fc_b])
print(scores.size()) # you should see [64, 10]
three_layer_convnet_test()

```

```
torch.Size([64, 10])
```

#### 5.0.4 Barebones PyTorch: Initialization

Let's write a couple utility methods to initialize the weight matrices for our models.

- `random_weight(shape)` initializes a weight tensor with the Kaiming normalization method.
- `zero_weight(shape)` initializes a weight tensor with all zeros. Useful for instantiating bias parameters.

The `random_weight` function uses the Kaiming normal initialization method, described in:

He et al, *Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification*, ICCV 2015, <https://arxiv.org/abs/1502.01852>

```

[49]: def random_weight(shape):
    """
    Create random Tensors for weights; setting requires_grad=True means that we
    want to compute gradients for these Tensors during the backward pass.
    We use Kaiming normalization: sqrt(2 / fan_in)
    """
    if len(shape) == 2: # FC weight
        fan_in = shape[0]
    else:
        fan_in = np.prod(shape[1:]) # conv weight [out_channel, in_channel, kH,
→kW]
    # randn is standard normal distribution generator.
    w = torch.randn(shape, device=device, dtype=dtype) * np.sqrt(2. / fan_in)
    w.requires_grad = True

```

```

    return w

def zero_weight(shape):
    return torch.zeros(shape, device=device, dtype=dtype, requires_grad=True)

# create a weight of shape [3 x 5]
# you should see the type `torch.cuda.FloatTensor` if you use GPU.
# Otherwise it should be `torch.FloatTensor`
random_weight((3, 5))

```

```

[49]: tensor([[ -1.6799e-01,  -3.5499e-01,   3.5931e-01,  -9.3700e-02,   7.2765e-02],
          [ -1.5015e-01,  -4.7320e-02,   7.8341e-01,  -5.4327e-01,   1.6318e+00],
          [  4.4994e-04,  -4.0029e-01,   2.4899e-01,  -2.2449e+00,   6.5217e-02]],
        device='cuda:0', requires_grad=True)

```

### 5.0.5 Barebones PyTorch: Check Accuracy

When training the model we will use the following function to check the accuracy of our model on the training or validation sets.

When checking accuracy we don't need to compute any gradients; as a result we don't need PyTorch to build a computational graph for us when we compute scores. To prevent a graph from being built we scope our computation under a `torch.no_grad()` context manager.

```

[50]: def check_accuracy_part2(loader, model_fn, params):
        """
        Check the accuracy of a classification model.

        Inputs:
        - loader: A DataLoader for the data split we want to check
        - model_fn: A function that performs the forward pass of the model,
                    with the signature scores = model_fn(x, params)
        - params: List of PyTorch Tensors giving parameters of the model

        Returns: Nothing, but prints the accuracy of the model
        """
        split = 'val' if loader.dataset.train else 'test'
        print('Checking accuracy on the %s set' % split)
        num_correct, num_samples = 0, 0
        with torch.no_grad():
            for x, y in loader:
                x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
                y = y.to(device=device, dtype=torch.int64)
                scores = model_fn(x, params)
                _, preds = scores.max(1)
                num_correct += (preds == y).sum()
                num_samples += preds.size(0)
            acc = float(num_correct) / num_samples

```

```
print('Got %d / %d correct (%.2f%%)' % (num_correct, num_samples, 100 *
→acc))
```

## 5.0.6 BareBones PyTorch: Training Loop

We can now set up a basic training loop to train our network. We will train the model using stochastic gradient descent without momentum. We will use `torch.functional.cross_entropy` to compute the loss; you can [read about it here](#).

The training loop takes as input the neural network function, a list of initialized parameters (`w1`, `w2` in our example), and learning rate.

```
[51]: def train_part2(model_fn, params, learning_rate):
      """
      Train a model on CIFAR-10.

      Inputs:
      - model_fn: A Python function that performs the forward pass of the model.
        It should have the signature scores = model_fn(x, params) where x is a
        PyTorch Tensor of image data, params is a list of PyTorch Tensors giving
        model weights, and scores is a PyTorch Tensor of shape (N, C) giving
        scores for the elements in x.
      - params: List of PyTorch Tensors giving weights for the model
      - learning_rate: Python scalar giving the learning rate to use for SGD

      Returns: Nothing
      """
      for t, (x, y) in enumerate(loader_train):
          # Move the data to the proper device (GPU or CPU)
          x = x.to(device=device, dtype=dtype)
          y = y.to(device=device, dtype=torch.long)

          # Forward pass: compute scores and loss
          scores = model_fn(x, params)
          loss = F.cross_entropy(scores, y)

          # Backward pass: PyTorch figures out which Tensors in the computational
          # graph has requires_grad=True and uses backpropagation to compute the
          # gradient of the loss with respect to these Tensors, and stores the
          # gradients in the .grad attribute of each Tensor.
          loss.backward()

          # Update parameters. We don't want to backpropagate through the
          # parameter updates, so we scope the updates under a torch.no_grad()
          # context manager to prevent a computational graph from being built.
          with torch.no_grad():
              for w in params:
                  w -= learning_rate * w.grad
```

```

        # Manually zero the gradients after running the backward pass
        w.grad.zero_()

    if t % print_every == 0:
        print('Iteration %d, loss = %.4f' % (t, loss.item()))
        check_accuracy_part2(loader_val, model_fn, params)
        print()

```

### 5.0.7 BareBones PyTorch: Train a Two-Layer Network

Now we are ready to run the training loop. We need to explicitly allocate tensors for the fully connected weights,  $w_1$  and  $w_2$ .

Each minibatch of CIFAR has 64 examples, so the tensor shape is  $[64, 3, 32, 32]$ .

After flattening,  $x$  shape should be  $[64, 3 * 32 * 32]$ . This will be the size of the first dimension of  $w_1$ . The second dimension of  $w_1$  is the hidden layer size, which will also be the first dimension of  $w_2$ .

Finally, the output of the network is a 10-dimensional vector that represents the probability distribution over 10 classes.

You don't need to tune any hyperparameters but you should see accuracies above 40% after training for one epoch.

```

[52]: hidden_layer_size = 4000
      learning_rate = 1e-2

      w1 = random_weight((3 * 32 * 32, hidden_layer_size))
      w2 = random_weight((hidden_layer_size, 10))

      train_part2(two_layer_fc, [w1, w2], learning_rate)

```

```

Iteration 0, loss = 3.6120
Checking accuracy on the val set
Got 136 / 1000 correct (13.60%)

```

```

Iteration 100, loss = 1.9961
Checking accuracy on the val set
Got 372 / 1000 correct (37.20%)

```

```

Iteration 200, loss = 2.4316
Checking accuracy on the val set
Got 404 / 1000 correct (40.40%)

```

```

Iteration 300, loss = 1.8590
Checking accuracy on the val set
Got 388 / 1000 correct (38.80%)

```

```

Iteration 400, loss = 1.4745
Checking accuracy on the val set

```

Got 389 / 1000 correct (38.90%)

Iteration 500, loss = 1.5857

Checking accuracy on the val set

Got 437 / 1000 correct (43.70%)

Iteration 600, loss = 1.8299

Checking accuracy on the val set

Got 458 / 1000 correct (45.80%)

Iteration 700, loss = 1.8206

Checking accuracy on the val set

Got 441 / 1000 correct (44.10%)

### 5.0.8 BareBones PyTorch: Training a ConvNet

In the below you should use the functions defined above to train a three-layer convolutional network on CIFAR. The network should have the following architecture:

1. Convolutional layer (with bias) with 32 5x5 filters, with zero-padding of 2
2. ReLU
3. Convolutional layer (with bias) with 16 3x3 filters, with zero-padding of 1
4. ReLU
5. Fully-connected layer (with bias) to compute scores for 10 classes

You should initialize your weight matrices using the `random_weight` function defined above, and you should initialize your bias vectors using the `zero_weight` function above.

You don't need to tune any hyperparameters, but if everything works correctly you should achieve an accuracy above 42% after one epoch.

```
[53]: learning_rate = 3e-3

channel_1 = 32
channel_2 = 16

conv_w1 = None
conv_b1 = None
conv_w2 = None
conv_b2 = None
fc_w = None
fc_b = None

#####
# TODO: Initialize the parameters of a three-layer ConvNet.
→#
#####
conv_w1 = random_weight((32,3,5,5))
conv_b1 = zero_weight((32))
```

```

conv_w2 = random_weight((16,32,3,3))
conv_b2 = zero_weight((16))
fc_w = random_weight((16*32*32,10))
fc_b = zero_weight((10))
pass
#####
#                                     END OF YOUR CODE                                     #
#→#
#####

params = [conv_w1, conv_b1, conv_w2, conv_b2, fc_w, fc_b]
train_part2(three_layer_convnet, params, learning_rate)

```

Iteration 0, loss = 4.1246  
Checking accuracy on the val set  
Got 151 / 1000 correct (15.10%)

Iteration 100, loss = 1.9132  
Checking accuracy on the val set  
Got 397 / 1000 correct (39.70%)

Iteration 200, loss = 1.8271  
Checking accuracy on the val set  
Got 440 / 1000 correct (44.00%)

Iteration 300, loss = 1.6976  
Checking accuracy on the val set  
Got 455 / 1000 correct (45.50%)

Iteration 400, loss = 1.4690  
Checking accuracy on the val set  
Got 453 / 1000 correct (45.30%)

Iteration 500, loss = 1.5157  
Checking accuracy on the val set  
Got 471 / 1000 correct (47.10%)

Iteration 600, loss = 1.4883  
Checking accuracy on the val set  
Got 479 / 1000 correct (47.90%)

Iteration 700, loss = 1.5758  
Checking accuracy on the val set  
Got 477 / 1000 correct (47.70%)

## 6 Part III. PyTorch Module API

Barebone PyTorch requires that we track all the parameter tensors by hand. This is fine for small networks with a few tensors, but it would be extremely inconvenient and error-prone to track tens or hundreds of tensors in larger networks.

PyTorch provides the `nn.Module` API for you to define arbitrary network architectures, while tracking every learnable parameters for you. In Part II, we implemented SGD ourselves. PyTorch also provides the `torch.optim` package that implements all the common optimizers, such as RMSProp, Adagrad, and Adam. It even supports approximate second-order methods like L-BFGS! You can refer to the [doc](#) for the exact specifications of each optimizer.

To use the Module API, follow the steps below:

1. Subclass `nn.Module`. Give your network class an intuitive name like `TwoLayerFC`.
2. In the constructor `__init__()`, define all the layers you need as class attributes. Layer objects like `nn.Linear` and `nn.Conv2d` are themselves `nn.Module` subclasses and contain learnable parameters, so that you don't have to instantiate the raw tensors yourself. `nn.Module` will track these internal parameters for you. Refer to the [doc](#) to learn more about the dozens of builtin layers. **Warning:** don't forget to call the `super().__init__()` first!
3. In the `forward()` method, define the *connectivity* of your network. You should use the attributes defined in `__init__` as function calls that take tensor as input and output the "transformed" tensor. Do *not* create any new layers with learnable parameters in `forward()`! All of them must be declared upfront in `__init__`.

After you define your Module subclass, you can instantiate it as an object and call it just like the NN forward function in part II.

### 6.0.1 Module API: Two-Layer Network

Here is a concrete example of a 2-layer fully connected network:

```
[54]: class TwoLayerFC(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super().__init__()
        # assign layer objects to class attributes
        self.fc1 = nn.Linear(input_size, hidden_size)
        # nn.init package contains convenient initialization methods
        # http://pytorch.org/docs/master/nn.html#torch-nn-init
        nn.init.kaiming_normal_(self.fc1.weight)
        self.fc2 = nn.Linear(hidden_size, num_classes)
        nn.init.kaiming_normal_(self.fc2.weight)

    def forward(self, x):
        # forward always defines connectivity
        x = flatten(x)
        scores = self.fc2(F.relu(self.fc1(x)))
        return scores

def test_TwoLayerFC():
```

```

    input_size = 50
    x = torch.zeros((64, input_size), dtype=dtype)  # minibatch size 64,
→feature dimension 50
    model = TwoLayerFC(input_size, 42, 10)
    scores = model(x)
    print(scores.size())  # you should see [64, 10]
test_TwoLayerFC()

```

```
torch.Size([64, 10])
```

## 6.0.2 Module API: Three-Layer ConvNet

It's your turn to implement a 3-layer ConvNet followed by a fully connected layer. The network architecture should be the same as in Part II:

1. Convolutional layer with channel\_1 5x5 filters with zero-padding of 2
2. ReLU
3. Convolutional layer with channel\_2 3x3 filters with zero-padding of 1
4. ReLU
5. Fully-connected layer to num\_classes classes

You should initialize the weight matrices of the model using the Kaiming normal initialization method.

**HINT:** <http://pytorch.org/docs/stable/nn.html#conv2d>

After you implement the three-layer ConvNet, the test\_ThreeLayerConvNet function will run your implementation; it should print (64, 10) for the shape of the output scores.

```

[55]: class ThreeLayerConvNet(nn.Module):
    def __init__(self, in_channel, channel_1, channel_2, num_classes):
        super(ThreeLayerConvNet, self).__init__()

→#####
    # TODO: Set up the layers you need for a three-layer ConvNet with the
→#
    # architecture defined above.
→#

→#####
    self.cnn1 = nn.Conv2d(in_channels=in_channel, out_channels=channel_1,
→kernel_size=(5,5),
                                padding=2)
    nn.init.kaiming_normal_(self.cnn1.weight)
    self.relu = nn.ReLU()
    self.cnn2 = nn.Conv2d(in_channels=channel_1, out_channels=channel_2,
→kernel_size=(3,3),
                                padding=1)
    nn.init.kaiming_normal_(self.cnn2.weight)
    self.fc = nn.Linear(channel_2*32*32, num_classes)

```



```

        nn.init.kaiming_normal_(self.fc.weight)
        pass

    #####
    #                                     END OF YOUR CODE
    #
    #####

    def forward(self, x):
        scores = None

    #####
    # TODO: Implement the forward function for a 3-layer ConvNet. you
    # should use the layers you defined in __init__ and specify the
    # connectivity of those layers in forward()
    #####
    x2 = self.relu(self.cnn2(self.relu(self.cnn1(x))))
    x2 = flatten(x2)
    scores = self.fc(x2)
    pass

    #####
    #                                     END OF YOUR CODE
    #
    #####

    return scores

def test_ThreeLayerConvNet():
    x = torch.zeros((64, 3, 32, 32), dtype=dtype) # minibatch size 64, image
    size [3, 32, 32]
    model = ThreeLayerConvNet(in_channel=3, channel_1=12, channel_2=8,
    num_classes=10)
    scores = model(x)
    print(scores.size()) # you should see [64, 10]
test_ThreeLayerConvNet()

```

```
torch.Size([64, 10])
```

### 6.0.3 Module API: Check Accuracy

Given the validation or test set, we can check the classification accuracy of a neural network.

This version is slightly different from the one in part II. You don't manually pass in the parameters anymore.

```
[56]: def check_accuracy_part34(loader, model):
    if loader.dataset.train:
        print('Checking accuracy on validation set')
    else:
        print('Checking accuracy on test set')
    num_correct = 0
    num_samples = 0
    model.eval() # set model to evaluation mode
    with torch.no_grad():
        for x, y in loader:
            x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
            y = y.to(device=device, dtype=torch.long)
            scores = model(x)
            _, preds = scores.max(1)
            num_correct += (preds == y).sum()
            num_samples += preds.size(0)
    acc = float(num_correct) / num_samples
    print('Got %d / %d correct (%.2f)' % (num_correct, num_samples, 100 *
    →acc))
```

### 6.0.4 Module API: Training Loop

We also use a slightly different training loop. Rather than updating the values of the weights ourselves, we use an Optimizer object from the torch.optim package, which abstract the notion of an optimization algorithm and provides implementations of most of the algorithms commonly used to optimize neural networks.

```
[57]: def train_part34(model, optimizer, epochs=1):
    """
    Train a model on CIFAR-10 using the PyTorch Module API.

    Inputs:
    - model: A PyTorch Module giving the model to train.
    - optimizer: An Optimizer object we will use to train the model
    - epochs: (Optional) A Python integer giving the number of epochs to train
    →for

    Returns: Nothing, but prints model accuracies during training.
    """
    model = model.to(device=device) # move the model parameters to CPU/GPU
    for e in range(epochs):
        for t, (x, y) in enumerate(loader_train):
            model.train() # put model to training mode
```

```

x = x.to(device=device, dtype=dtype) # move to device, e.g. GPU
y = y.to(device=device, dtype=torch.long)

scores = model(x)
loss = F.cross_entropy(scores, y)

# Zero out all of the gradients for the variables which the
→optimizer
# will update.
optimizer.zero_grad()

# This is the backwards pass: compute the gradient of the loss with
# respect to each parameter of the model.
loss.backward()

# Actually update the parameters of the model using the gradients
# computed by the backwards pass.
optimizer.step()

if t % print_every == 0:
    print('Iteration %d, loss = %.4f' % (t, loss.item()))
    check_accuracy_part34(loader_val, model)
    print()

```

### 6.0.5 Module API: Train a Two-Layer Network

Now we are ready to run the training loop. In contrast to part II, we don't explicitly allocate parameter tensors anymore.

Simply pass the input size, hidden layer size, and number of classes (i.e. output size) to the constructor of `TwoLayerFC`.

You also need to define an optimizer that tracks all the learnable parameters inside `TwoLayerFC`.

You don't need to tune any hyperparameters, but you should see model accuracies above 40% after training for one epoch.

```

[58]: hidden_layer_size = 4000
learning_rate = 1e-2
model = TwoLayerFC(3 * 32 * 32, hidden_layer_size, 10)
optimizer = optim.SGD(model.parameters(), lr=learning_rate)

train_part34(model, optimizer)

```

```

Iteration 0, loss = 3.4846
Checking accuracy on validation set
Got 142 / 1000 correct (14.20)

```

```

Iteration 100, loss = 2.2141
Checking accuracy on validation set

```

```

Got 337 / 1000 correct (33.70)

Iteration 200, loss = 1.9447
Checking accuracy on validation set
Got 391 / 1000 correct (39.10)

Iteration 300, loss = 2.2036
Checking accuracy on validation set
Got 387 / 1000 correct (38.70)

Iteration 400, loss = 1.6199
Checking accuracy on validation set
Got 371 / 1000 correct (37.10)

Iteration 500, loss = 1.6941
Checking accuracy on validation set
Got 380 / 1000 correct (38.00)

Iteration 600, loss = 1.8460
Checking accuracy on validation set
Got 433 / 1000 correct (43.30)

Iteration 700, loss = 1.7330
Checking accuracy on validation set
Got 451 / 1000 correct (45.10)

```

### 6.0.6 Module API: Train a Three-Layer ConvNet

You should now use the Module API to train a three-layer ConvNet on CIFAR. This should look very similar to training the two-layer network! You don't need to tune any hyperparameters, but you should achieve above 45% after training for one epoch.

You should train the model using stochastic gradient descent without momentum.

```

[59]: learning_rate = 3e-3
      channel_1 = 32
      channel_2 = 16

      model = None
      optimizer = None
      #####
      # TODO: Instantiate your ThreeLayerConvNet model and a corresponding optimizer_
      ↪ #
      #####
      model = ThreeLayerConvNet(3, channel_1, channel_2, 10)
      optimizer = optim.SGD(model.parameters(), lr=learning_rate)
      pass
      #####

```

```
# END OF YOUR CODE
#####

train_part34(model, optimizer)
```

```
Iteration 0, loss = 3.2429
Checking accuracy on validation set
Got 120 / 1000 correct (12.00)
```

```
Iteration 100, loss = 1.7959
Checking accuracy on validation set
Got 333 / 1000 correct (33.30)
```

```
Iteration 200, loss = 1.7593
Checking accuracy on validation set
Got 417 / 1000 correct (41.70)
```

```
Iteration 300, loss = 1.5599
Checking accuracy on validation set
Got 435 / 1000 correct (43.50)
```

```
Iteration 400, loss = 1.6712
Checking accuracy on validation set
Got 450 / 1000 correct (45.00)
```

```
Iteration 500, loss = 1.3839
Checking accuracy on validation set
Got 466 / 1000 correct (46.60)
```

```
Iteration 600, loss = 1.3204
Checking accuracy on validation set
Got 473 / 1000 correct (47.30)
```

```
Iteration 700, loss = 1.5330
Checking accuracy on validation set
Got 478 / 1000 correct (47.80)
```

## 7 Part IV. PyTorch Sequential API

Part III introduced the PyTorch Module API, which allows you to define arbitrary learnable layers and their connectivity.

For simple models like a stack of feed forward layers, you still need to go through 3 steps: subclass `nn.Module`, assign layers to class attributes in `__init__`, and call each layer one by one in `forward()`. Is there a more convenient way?

Fortunately, PyTorch provides a container Module called `nn.Sequential`, which merges the above steps into one. It is not as flexible as `nn.Module`, because you cannot specify more complex

topology than a feed-forward stack, but it's good enough for many use cases.

### 7.0.1 Sequential API: Two-Layer Network

Let's see how to rewrite our two-layer fully connected network example with `nn.Sequential`, and train it using the training loop defined above.

Again, you don't need to tune any hyperparameters here, but you should achieve above 40% accuracy after one epoch of training.

```
[60]: # We need to wrap `flatten` function in a module in order to stack it
      # in nn.Sequential
      class Flatten(nn.Module):
          def forward(self, x):
              return flatten(x)

      hidden_layer_size = 4000
      learning_rate = 1e-2

      model = nn.Sequential(
          Flatten(),
          nn.Linear(3 * 32 * 32, hidden_layer_size),
          nn.ReLU(),
          nn.Linear(hidden_layer_size, 10),
      )

      # you can use Nesterov momentum in optim.SGD
      optimizer = optim.SGD(model.parameters(), lr=learning_rate,
                             momentum=0.9, nesterov=True)

      train_part34(model, optimizer)
```

```
Iteration 0, loss = 2.3848
Checking accuracy on validation set
Got 163 / 1000 correct (16.30)
```

```
Iteration 100, loss = 1.7335
Checking accuracy on validation set
Got 407 / 1000 correct (40.70)
```

```
Iteration 200, loss = 1.7850
Checking accuracy on validation set
Got 427 / 1000 correct (42.70)
```

```
Iteration 300, loss = 1.5170
Checking accuracy on validation set
Got 424 / 1000 correct (42.40)
```

```
Iteration 400, loss = 1.9202
```

```
Checking accuracy on validation set
Got 409 / 1000 correct (40.90)
```

```
Iteration 500, loss = 1.5282
Checking accuracy on validation set
Got 436 / 1000 correct (43.60)
```

```
Iteration 600, loss = 1.4944
Checking accuracy on validation set
Got 453 / 1000 correct (45.30)
```

```
Iteration 700, loss = 1.7569
Checking accuracy on validation set
Got 444 / 1000 correct (44.40)
```

## 7.0.2 Sequential API: Three-Layer ConvNet

Here you should use `nn.Sequential` to define and train a three-layer ConvNet with the same architecture we used in Part III:

1. Convolutional layer (with bias) with 32 5x5 filters, with zero-padding of 2
2. ReLU
3. Convolutional layer (with bias) with 16 3x3 filters, with zero-padding of 1
4. ReLU
5. Fully-connected layer (with bias) to compute scores for 10 classes

You should initialize your weight matrices using the `random_weight` function defined above, and you should initialize your bias vectors using the `zero_weight` function above.

You should optimize your model using stochastic gradient descent with Nesterov momentum 0.9.

Again, you don't need to tune any hyperparameters but you should see accuracy above 55% after one epoch of training.

```
[61]: channel_1 = 32
      channel_2 = 16
      learning_rate = 1e-2

      model = None
      optimizer = None

      #####
      # TODO: Rewrite the 2-layer ConvNet with bias from Part III with the
      # Sequential API.
      #####
      model = nn.Sequential(
          nn.Conv2d(in_channels=3,out_channels=channel_1,kernel_size=(5,5),padding=2),
```

```

    nn.ReLU(),
    nn.
    →Conv2d(in_channels=channel_1,out_channels=channel_2,kernel_size=(3,3),padding=1),
    nn.ReLU(),
    nn.Flatten(),
    nn.Linear(channel_2*32*32,10)
)

optimizer = optim.SGD(model.parameters(), lr=learning_rate,
                       momentum=0.9, nesterov=True)

pass
#####
#                                     END OF YOUR CODE
#####

train_part34(model, optimizer)

```

Iteration 0, loss = 2.3067  
 Checking accuracy on validation set  
 Got 184 / 1000 correct (18.40)

Iteration 100, loss = 1.4464  
 Checking accuracy on validation set  
 Got 442 / 1000 correct (44.20)

Iteration 200, loss = 1.8018  
 Checking accuracy on validation set  
 Got 446 / 1000 correct (44.60)

Iteration 300, loss = 1.4289  
 Checking accuracy on validation set  
 Got 445 / 1000 correct (44.50)

Iteration 400, loss = 1.4361  
 Checking accuracy on validation set  
 Got 537 / 1000 correct (53.70)

Iteration 500, loss = 1.2079  
 Checking accuracy on validation set  
 Got 525 / 1000 correct (52.50)

Iteration 600, loss = 1.5931  
 Checking accuracy on validation set  
 Got 491 / 1000 correct (49.10)

Iteration 700, loss = 1.4455  
 Checking accuracy on validation set



Got 557 / 1000 correct (55.70)

## 8 Part V. CIFAR-10 open-ended challenge

In this section, you can experiment with whatever ConvNet architecture you'd like on CIFAR-10.

Now it's your job to experiment with architectures, hyperparameters, loss functions, and optimizers to train a model that achieves **at least 70%** accuracy on the CIFAR-10 **validation** set within 10 epochs. You can use the `check_accuracy` and `train` functions from above. You can use either `nn.Module` or `nn.Sequential` API.

Describe what you did at the end of this notebook.

Here are the official API documentation for each component. One note: what we call in the class "spatial batch norm" is called "BatchNorm2D" in PyTorch.

- Layers in torch.nn package: <http://pytorch.org/docs/stable/nn.html>
- Activations: <http://pytorch.org/docs/stable/nn.html#non-linear-activations>
- Loss functions: <http://pytorch.org/docs/stable/nn.html#loss-functions>
- Optimizers: <http://pytorch.org/docs/stable/optim.html>

### 8.0.1 Things you might try:

- **Filter size:** Above we used 5x5; would smaller filters be more efficient?
- **Number of filters:** Above we used 32 filters. Do more or fewer do better?
- **Pooling vs Strided Convolution:** Do you use max pooling or just stride convolutions?
- **Batch normalization:** Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- **Network architecture:** The network above has two layers of trainable parameters. Can you do better with a deep network? Good architectures to try include:
  - [conv-relu-pool]xN -> [affine]xM -> [softmax or SVM]
  - [conv-relu-conv-relu-pool]xN -> [affine]xM -> [softmax or SVM]
  - [batchnorm-relu-conv]xN -> [affine]xM -> [softmax or SVM]
- **Global Average Pooling:** Instead of flattening and then having multiple affine layers, perform convolutions until your image gets small (7x7 or so) and then perform an average pooling operation to get to a 1x1 image picture (1, 1, Filter#), which is then reshaped into a (Filter#) vector. This is used in [Google's Inception Network](#) (See Table 1 for their architecture).
- **Regularization:** Add l2 weight regularization, or perhaps use Dropout.

### 8.0.2 Tips for training

For each network architecture that you try, you should tune the learning rate and other hyperparameters. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations

- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.
- You should use the validation set for hyperparameter search, and save your test set for evaluating your architecture on the best parameters as selected by the validation set.

### 8.0.3 Going above and beyond

If you are feeling adventurous there are many other features you can implement to try and improve your performance. You are **not required** to implement any of these, but don't miss the fun if you have time!

- Alternative optimizers: you can try Adam, Adagrad, RMSprop, etc.
- Alternative activation functions such as leaky ReLU, parametric ReLU, ELU, or MaxOut.
- Model ensembles
- Data augmentation
- New Architectures
- [ResNets](#) where the input from the previous layer is added to the output.
- [DenseNets](#) where inputs into previous layers are concatenated together.
- [This blog has an in-depth overview](#)

### 8.0.4 Have fun and happy training!

```
[70]: #####
# TODO:
#
# Experiment with any architectures, optimizers, and hyperparameters.
#
# Achieve AT LEAST 70% accuracy on the *validation set* within 10 epochs.
#
#
# Note that you can use the check_accuracy function to evaluate on either
# the test set or the validation set, by passing either loader_test or
# loader_val as the second argument to check_accuracy. You should not touch
# the test set until you have finished your architecture and hyperparameter
# tuning, and only run the test set once at the end to report a final value.
#####
model = None
optimizer = None
```

```

learning_rate = 1e-3
model = nn.Sequential(
    nn.Conv2d(in_channels=3,out_channels=16,kernel_size=(3,3),padding=2),
    nn.Conv2d(in_channels=16,out_channels=16,kernel_size=(3,3),padding=2),
    nn.GroupNorm(4,16),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2),
    nn.Dropout(0.4),
    nn.Conv2d(in_channels=16,out_channels=32,kernel_size=(3,3),padding=1),
    nn.Conv2d(in_channels=32,out_channels=32,kernel_size=(3,3),padding=1),
    nn.Conv2d(in_channels=32,out_channels=32,kernel_size=(3,3),padding=1),
    nn.GroupNorm(8,32),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2),
    nn.Dropout(0.4),
    nn.Conv2d(in_channels=32,out_channels=64,kernel_size=(3,3),padding=1),
    nn.Conv2d(in_channels=64,out_channels=64,kernel_size=(3,3),padding=1),
    nn.Conv2d(in_channels=64,out_channels=64,kernel_size=(3,3),padding=1),
    nn.GroupNorm(16,64),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2),
    nn.Dropout(0.4),
    nn.Flatten(),
    nn.Linear(64*4*4,10)
)

optimizer = optim.Adam(model.parameters(), lr=learning_rate)
pass
#####
#                               END OF YOUR CODE
#####

# You should get at least 70% accuracy
train_part34(model, optimizer, epochs=10)

```

```

Iteration 0, loss = 2.4730
Checking accuracy on validation set
Got 117 / 1000 correct (11.70)

```

```

Iteration 100, loss = 2.1408
Checking accuracy on validation set
Got 297 / 1000 correct (29.70)

```

```

Iteration 200, loss = 1.6165
Checking accuracy on validation set
Got 435 / 1000 correct (43.50)

```

Iteration 300, loss = 1.6952  
Checking accuracy on validation set  
Got 443 / 1000 correct (44.30)

Iteration 400, loss = 1.5159  
Checking accuracy on validation set  
Got 453 / 1000 correct (45.30)

Iteration 500, loss = 1.4734  
Checking accuracy on validation set  
Got 479 / 1000 correct (47.90)

Iteration 600, loss = 1.5004  
Checking accuracy on validation set  
Got 485 / 1000 correct (48.50)

Iteration 700, loss = 1.5061  
Checking accuracy on validation set  
Got 506 / 1000 correct (50.60)

Iteration 0, loss = 1.5558  
Checking accuracy on validation set  
Got 552 / 1000 correct (55.20)

Iteration 100, loss = 1.3707  
Checking accuracy on validation set  
Got 546 / 1000 correct (54.60)

Iteration 200, loss = 1.5531  
Checking accuracy on validation set  
Got 561 / 1000 correct (56.10)

Iteration 300, loss = 1.1997  
Checking accuracy on validation set  
Got 560 / 1000 correct (56.00)

Iteration 400, loss = 1.1147  
Checking accuracy on validation set  
Got 600 / 1000 correct (60.00)

Iteration 500, loss = 1.2809  
Checking accuracy on validation set  
Got 589 / 1000 correct (58.90)

Iteration 600, loss = 1.2585  
Checking accuracy on validation set  
Got 606 / 1000 correct (60.60)

Iteration 700, loss = 1.1818  
Checking accuracy on validation set  
Got 613 / 1000 correct (61.30)

Iteration 0, loss = 1.2039  
Checking accuracy on validation set  
Got 631 / 1000 correct (63.10)

Iteration 100, loss = 1.3887  
Checking accuracy on validation set  
Got 630 / 1000 correct (63.00)

Iteration 200, loss = 1.0475  
Checking accuracy on validation set  
Got 626 / 1000 correct (62.60)

Iteration 300, loss = 1.1899  
Checking accuracy on validation set  
Got 637 / 1000 correct (63.70)

Iteration 400, loss = 1.2673  
Checking accuracy on validation set  
Got 635 / 1000 correct (63.50)

Iteration 500, loss = 1.1044  
Checking accuracy on validation set  
Got 651 / 1000 correct (65.10)

Iteration 600, loss = 1.3838  
Checking accuracy on validation set  
Got 649 / 1000 correct (64.90)

Iteration 700, loss = 1.0814  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Iteration 0, loss = 0.9574  
Checking accuracy on validation set  
Got 653 / 1000 correct (65.30)

Iteration 100, loss = 1.0782  
Checking accuracy on validation set  
Got 649 / 1000 correct (64.90)

Iteration 200, loss = 1.1186  
Checking accuracy on validation set  
Got 657 / 1000 correct (65.70)

Iteration 300, loss = 1.0609  
Checking accuracy on validation set  
Got 680 / 1000 correct (68.00)

Iteration 400, loss = 1.0633  
Checking accuracy on validation set  
Got 658 / 1000 correct (65.80)

Iteration 500, loss = 1.1267  
Checking accuracy on validation set  
Got 649 / 1000 correct (64.90)

Iteration 600, loss = 0.9903  
Checking accuracy on validation set  
Got 664 / 1000 correct (66.40)

Iteration 700, loss = 0.9965  
Checking accuracy on validation set  
Got 656 / 1000 correct (65.60)

Iteration 0, loss = 1.1092  
Checking accuracy on validation set  
Got 679 / 1000 correct (67.90)

Iteration 100, loss = 1.0461  
Checking accuracy on validation set  
Got 671 / 1000 correct (67.10)

Iteration 200, loss = 0.9576  
Checking accuracy on validation set  
Got 681 / 1000 correct (68.10)

Iteration 300, loss = 1.0632  
Checking accuracy on validation set  
Got 676 / 1000 correct (67.60)

Iteration 400, loss = 1.0333  
Checking accuracy on validation set  
Got 674 / 1000 correct (67.40)

Iteration 500, loss = 1.0733  
Checking accuracy on validation set  
Got 687 / 1000 correct (68.70)

Iteration 600, loss = 1.1887  
Checking accuracy on validation set  
Got 693 / 1000 correct (69.30)

Iteration 700, loss = 0.8796  
Checking accuracy on validation set  
Got 688 / 1000 correct (68.80)

Iteration 0, loss = 1.0895  
Checking accuracy on validation set  
Got 667 / 1000 correct (66.70)

Iteration 100, loss = 1.0018  
Checking accuracy on validation set  
Got 697 / 1000 correct (69.70)

Iteration 200, loss = 0.9658  
Checking accuracy on validation set  
Got 688 / 1000 correct (68.80)

Iteration 300, loss = 1.1203  
Checking accuracy on validation set  
Got 674 / 1000 correct (67.40)

Iteration 400, loss = 1.0450  
Checking accuracy on validation set  
Got 690 / 1000 correct (69.00)

Iteration 500, loss = 1.0032  
Checking accuracy on validation set  
Got 700 / 1000 correct (70.00)

Iteration 600, loss = 1.0238  
Checking accuracy on validation set  
Got 686 / 1000 correct (68.60)

Iteration 700, loss = 1.0911  
Checking accuracy on validation set  
Got 694 / 1000 correct (69.40)

Iteration 0, loss = 1.1055  
Checking accuracy on validation set  
Got 701 / 1000 correct (70.10)

Iteration 100, loss = 1.1320  
Checking accuracy on validation set  
Got 701 / 1000 correct (70.10)

Iteration 200, loss = 1.0483  
Checking accuracy on validation set  
Got 694 / 1000 correct (69.40)

Iteration 300, loss = 1.0196  
Checking accuracy on validation set  
Got 698 / 1000 correct (69.80)

Iteration 400, loss = 0.9929  
Checking accuracy on validation set  
Got 718 / 1000 correct (71.80)

Iteration 500, loss = 0.9545  
Checking accuracy on validation set  
Got 710 / 1000 correct (71.00)

Iteration 600, loss = 0.7093  
Checking accuracy on validation set  
Got 695 / 1000 correct (69.50)

Iteration 700, loss = 0.8338  
Checking accuracy on validation set  
Got 715 / 1000 correct (71.50)

Iteration 0, loss = 0.8267  
Checking accuracy on validation set  
Got 711 / 1000 correct (71.10)

Iteration 100, loss = 0.8627  
Checking accuracy on validation set  
Got 722 / 1000 correct (72.20)

Iteration 200, loss = 0.8839  
Checking accuracy on validation set  
Got 711 / 1000 correct (71.10)

Iteration 300, loss = 1.2598  
Checking accuracy on validation set  
Got 713 / 1000 correct (71.30)

Iteration 400, loss = 0.8186  
Checking accuracy on validation set  
Got 720 / 1000 correct (72.00)

Iteration 500, loss = 0.8182  
Checking accuracy on validation set  
Got 710 / 1000 correct (71.00)

Iteration 600, loss = 1.0270  
Checking accuracy on validation set  
Got 720 / 1000 correct (72.00)



Iteration 700, loss = 0.8936  
Checking accuracy on validation set  
Got 714 / 1000 correct (71.40)

Iteration 0, loss = 1.0181  
Checking accuracy on validation set  
Got 717 / 1000 correct (71.70)

Iteration 100, loss = 0.8258  
Checking accuracy on validation set  
Got 716 / 1000 correct (71.60)

Iteration 200, loss = 0.9993  
Checking accuracy on validation set  
Got 728 / 1000 correct (72.80)

Iteration 300, loss = 1.1637  
Checking accuracy on validation set  
Got 707 / 1000 correct (70.70)

Iteration 400, loss = 0.9116  
Checking accuracy on validation set  
Got 733 / 1000 correct (73.30)

Iteration 500, loss = 1.0024  
Checking accuracy on validation set  
Got 714 / 1000 correct (71.40)

Iteration 600, loss = 0.9839  
Checking accuracy on validation set  
Got 723 / 1000 correct (72.30)

Iteration 700, loss = 1.0408  
Checking accuracy on validation set  
Got 711 / 1000 correct (71.10)

Iteration 0, loss = 0.8843  
Checking accuracy on validation set  
Got 731 / 1000 correct (73.10)

Iteration 100, loss = 0.9879  
Checking accuracy on validation set  
Got 733 / 1000 correct (73.30)

Iteration 200, loss = 0.9994  
Checking accuracy on validation set  
Got 721 / 1000 correct (72.10)

```
Iteration 300, loss = 0.7924
Checking accuracy on validation set
Got 728 / 1000 correct (72.80)
```

```
Iteration 400, loss = 0.8144
Checking accuracy on validation set
Got 713 / 1000 correct (71.30)
```

```
Iteration 500, loss = 1.2459
Checking accuracy on validation set
Got 729 / 1000 correct (72.90)
```

```
Iteration 600, loss = 0.8843
Checking accuracy on validation set
Got 731 / 1000 correct (73.10)
```

```
Iteration 700, loss = 0.9546
Checking accuracy on validation set
Got 729 / 1000 correct (72.90)
```

## 8.1 Describe what you did

In the cell below you should write an explanation of what you did, any additional features that you implemented, and/or any graphs that you made in the process of training and evaluating your network.

TODO:

I modified the architecture used for our trial three layer network by adding additional convolution layers, group norm layers, dropouts and max pooling layers. I used Adam optimizer with learning rate of 1e-3. Below is the architecture overview -

[Input Layer] -> [Conv-Conv-Group-ReLU-MaxPool-Dropout] -> [Conv-Conv-Conv-Group-ReLU-MaxPool-Dropout]x2 -> [Affine] -> [Softmax]

The final validation accuracy I achieved using this architecture is 72.9%

## 8.2 Test set -- run this only once

Now that we've gotten a result we're happy with, we test our final model on the test set (which you should store in `best_model`). Think about how this compares to your validation set accuracy.

```
[71]: best_model = model
      check_accuracy_part34(loader_test, best_model)
```

```
Checking accuracy on test set
Got 7158 / 10000 correct (71.58)
```